

# Machine Learning Homework 03

Author: Jesus Arturo Sol Navarro

## 1 Problem 6

Compute the first and second derivative of this likelihood w.r.t.  $\theta$ . Then compute first and second derivative of the log-likelihood  $\log \theta^t (1 - \theta)^h$ .

$$f(\theta) = \theta^t (1 - \theta)^h$$

$$\frac{df}{d\theta} = \frac{d}{d\theta}(\theta^t)(1 - \theta)^h + \theta^t \frac{d}{d\theta}(1 - \theta)^h = t\theta^{t-1}(1 - \theta)^h - \theta^t h(1 - \theta)^{h-1}$$

$$\frac{d^2 f}{d\theta^2} = \frac{d}{d\theta}(t\theta^{t-1})(1 - \theta)^h + t\theta^{t-1} \frac{d}{d\theta}(1 - \theta)^h + \frac{d}{d\theta}(h\theta^t)(1 - \theta)^{h-1} + h\theta^t \frac{d}{d\theta}(1 - \theta)^{h-1}$$

$$= t(t-1)\theta^{t-2}(1 - \theta)^h - t\theta^{t-1}h(1 - \theta)^{h-1} - ht\theta^{t-1}(1 - \theta)^{h-1} + h\theta^t(h-1)(1 - \theta)^{h-2}$$

$$\frac{d \log(f)}{d\theta} = \frac{d}{d\theta}(t \log(\theta) + h \log(1 - \theta)) = \frac{t}{\theta} - \frac{h}{1 - \theta}$$

$$\frac{d^2 \log(f)}{d\theta^2} = -\frac{t}{\theta^2} - \frac{h}{(1 - \theta)^2}$$

## 2 Problem 7

Show that for any differentiable, positive function  $f(\theta)$  every local maximum of  $\log f(\theta)$  is also a local maximum of  $f(\theta)$ . Considering this and the previous exercise, what is your conclusion?

Taking in consideration that  $f(\theta) > 0$  for all  $\theta$ , then  $h(\theta) = \log(f(\theta))$  is a well-defined, differentiable function.

Computing the first derivative and setting it to zero.

$$\frac{dh(\theta)}{d\theta} = \frac{d \log f(\theta)}{d\theta} = \frac{f'(\theta)}{f(\theta)} \stackrel{!}{=} 0$$

It can be identified that  $h'(\theta) = 0$  if and only if  $f'(\theta) = 0$ . Since the first derivative of both functions share the same zero points,  $h(\theta) = \log(f(\theta))$  has local maxima at the same points as  $f(\theta)$ .

Therefore, maximizing  $\log(f(\theta))$  is equivalent to maximizing  $f(\theta)$ , this simplifies greatly the process of finding MLE or MAP.

## 3 Problem 8

For the random variables:

$$\Theta \sim \text{Beta}(a, b) = \Theta^{a-1}(1 - \Theta)^{b-1}$$

$$X \sim \text{Binom}(N, \Theta) = \Theta^m(1 - \Theta)^{N-m}$$

Identification of the posterior distribution:

$$p(\Theta|X) = \frac{p(X|\Theta) \cdot p(\Theta)}{p(X)} \propto p(X|\Theta) \cdot p(\Theta) \propto \Theta^m(1 - \Theta)^{N-m} \cdot \Theta^{a-1}(1 - \Theta)^{b-1} \propto \Theta^{m+a-1}(1 - \Theta)^{N-m+b-1}$$

because of the equation  $N - m = l$  we get:

$$\Theta^{m+a-1}(1 - \Theta)^{l+b-1}$$

which is a Beta distribution for  $\Theta|X \sim \text{Beta}(m + a, l + b)$  The posterior mean of  $\Theta$  is:

$$\mathbf{E}[\Theta|D] = \frac{m + a}{m + a + l + b} = \frac{m}{m + a + l + b} + \frac{a}{m + a + l + b} = \frac{m}{m + l} \cdot \frac{m + l}{m + a + l + b} + \frac{a}{a + b} \cdot \frac{a + b}{m + a + l + b}$$

$$\frac{m+l}{m+l+a+b} = 1 - \lambda \text{ and } \frac{a+b}{m+a+l+b} = \lambda$$

So  $\frac{m}{m+l}$  is the maximum likelihood estimate and  $\frac{a}{a+b}$  is the prior mean value of  $\Theta$

## 4 Problem 9

$$\begin{aligned}
 \lambda_{MAP} &= \arg \max_{\lambda} p(\lambda|x, a, b) = \arg \max_{\lambda} \log p(\lambda|x, a, b) = \arg \max_{\lambda} (\log p(x|\lambda) + \log p(\lambda|a, b)) = \\
 &= \arg \max_{\lambda} (\log(\frac{\lambda^x \exp(-\lambda)}{x!}) + \log(\frac{b^a}{\Gamma(a)} \lambda^{a-1} \exp(-b\lambda))) = \\
 &= \arg \max_{\lambda} (x \log(\lambda) - \lambda + \log(\frac{1}{x!}) + \log(\frac{b^a}{\Gamma(a)}) + (a-1) \log(\lambda) - b\lambda) = \\
 &= \arg \max_{\lambda} ((x+a-1) \log(\lambda) - (1+b)\lambda) \\
 &\quad \frac{\delta((x+a-1) \log(\lambda) - (1+b)\lambda)}{\delta(\lambda)} = 0 \\
 &\quad \frac{x+a-1}{\lambda} - 1 - b = 0 \\
 &\quad \lambda_{MAP} = \frac{x+a-1}{1+b}
 \end{aligned}$$

## 5 Problem 10

### 5.1 Simulating data

```

def simulate_data(num_samples, tails_proba):
    """Simulate a sequence of i.i.d. coin flips.

    Tails are denoted as 1 and heads are denoted as 0.

    Parameters
    -----
    num_samples : int
        Number of samples to generate.
    tails_proba : float in range (0, 1)
        Probability of observing tails.

    Returns
    -----
    samples : array, shape (num_samples)
        Outcomes of simulated coin flips. Tails is 1 and heads is 0.
    """
    return np.random.choice([0, 1], size=(num_samples), p=[1 - tails_proba, tails_proba])

np.random.seed(123) # for reproducibility
num_samples = 20
tails_proba = 0.7
samples = simulate_data(num_samples, tails_proba)
print(samples)

```

## 5.2 Compute $\log p(\mathcal{D} \mid \theta)$ for different values of $\theta$

```
def compute_log_likelihood(theta, samples):
    """Compute log p(D | theta) for the given values of theta.

    Parameters
    -----
    theta : array, shape (num_points)
        Values of theta for which it's necessary to evaluate the log-likelihood.
    samples : array, shape (num_samples)
        Outcomes of simulated coin flips. Tails is 1 and heads is 0.

    Returns
    -----
    log_likelihood : array, shape (num_points)
        Values of log-likelihood for each value in theta.
    """

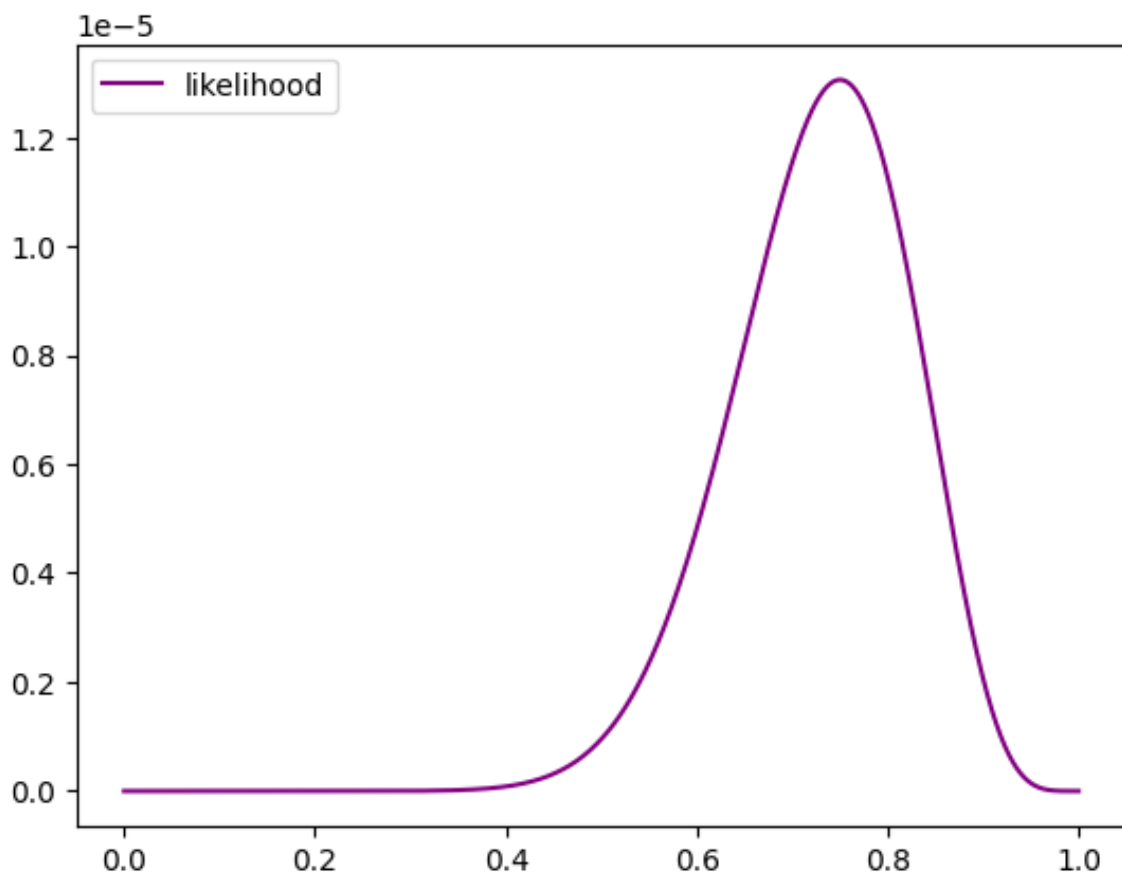
    number_of_tails = np.count_nonzero(samples==1)
    number_of_heads = len(samples) - number_of_tails
    likelihood = np.power(theta,number_of_tails) * (np.power((1-theta),number_of_heads))

    log_likelihood1 = number_of_tails*(np.log(theta))
    log_likelihood2 = number_of_heads*(np.log(1-theta))
    log_likelihood = log_likelihood1+log_likelihood2

    return log_likelihood

x = np.linspace(1e-5, 1-1e-5, 1000)
log_likelihood = compute_log_likelihood(x, samples)
likelihood = np.exp(log_likelihood)
plt.plot(x, likelihood, label='likelihood', c='purple')
plt.legend()

int_likelihood = 1.0 * np.mean(likelihood)
print(f'Integral = {int_likelihood:.4}')
```



### 5.3 Compute $p(\theta | a, b)$ for different values of $\theta$

```
def compute_log_prior(theta, a, b):
    """Compute log p(theta | a, b) for the given values of theta.

    Parameters
    -----
    theta : array, shape (num_points)
        Values of theta for which it's necessary to evaluate the log-prior.
    a, b: float
        Parameters of the prior Beta distribution.

    Returns
    -----
    log_prior : array, shape (num_points)
        Values of log-prior for each value in theta.

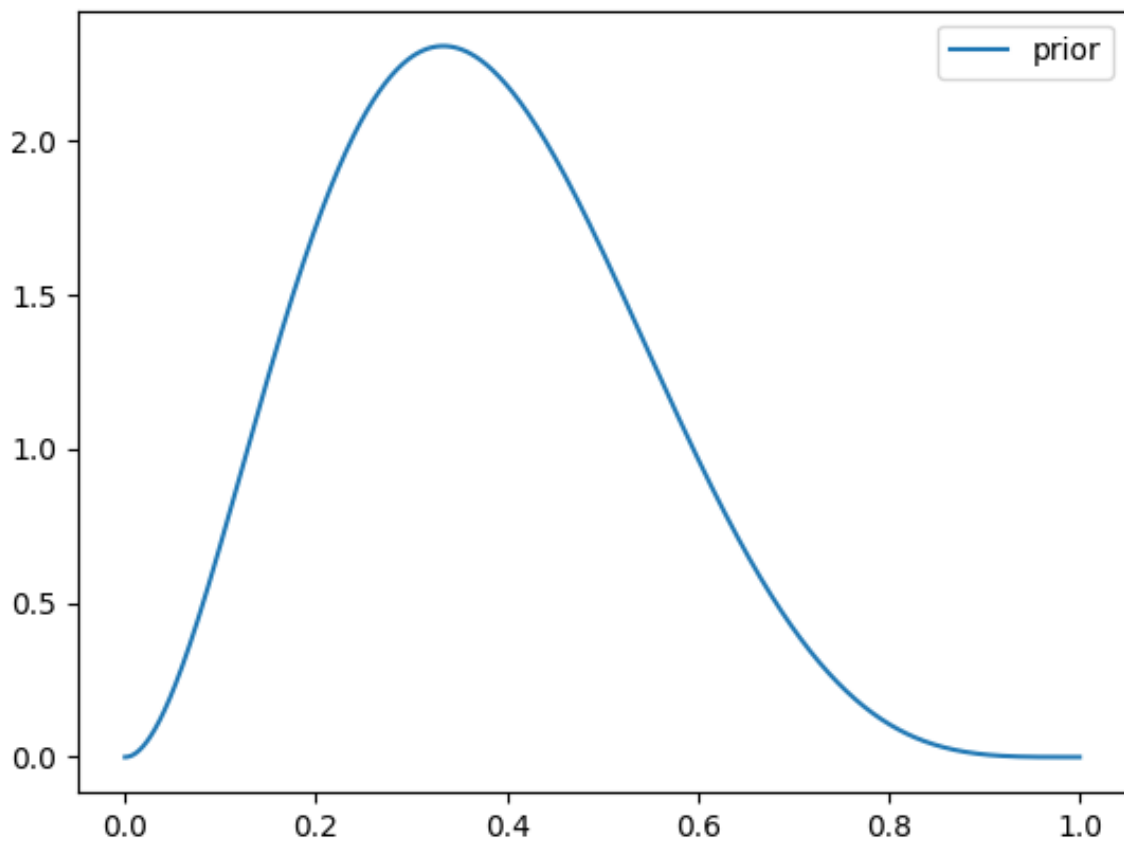
    """
    m= (a-1)*np.log(theta)
    n= (b-1)*np.log(1-theta)
    o= loggamma(a)+loggamma(b)-loggamma(a+b)
    logbeta = m+n-o

    return logbeta

x = np.linspace(1e-5, 1-1e-5, 1000)
a, b = 3, 5

# Plot the prior distribution
log_prior = compute_log_prior(x, a, b)
prior = np.exp(log_prior)
plt.plot(x, prior, label='prior')
plt.legend()

int_prior = 1.0 * np.mean(prior)
print(f'Integral = {int_prior:.4}')
```



## 5.4 Compute $\log p(\theta \mid \mathcal{D}, a, b)$ for different values of $\theta$

```
def compute_log_posterior(theta, samples, a, b):
    """Compute  $\log p(\theta \mid \mathcal{D}, a, b)$  for the given values of theta.

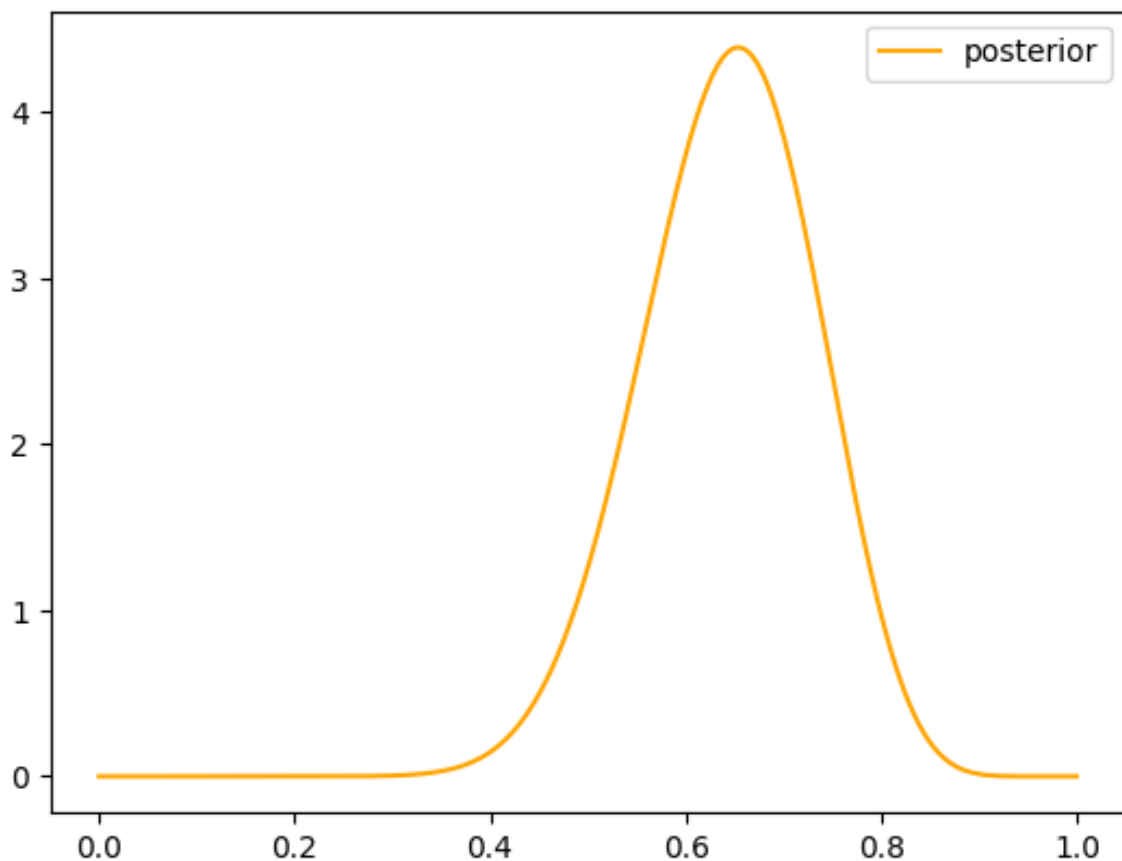
    Parameters
    -----
    theta : array, shape (num_points)
        Values of theta for which it's necessary to evaluate the log-prior.
    samples : array, shape (num_samples)
        Outcomes of simulated coin flips. Tails is 1 and heads is 0.
    a, b: float
        Parameters of the prior Beta distribution.

    Returns
    -----
    log_posterior : array, shape (num_points)
        Values of log-posterior for each value in theta.
    """
    number_of_tails = np.count_nonzero(samples==1)
    number_of_heads = len(samples) - number_of_tails
    log_posterior=compute_log_prior(theta, a+number_of_tails, b+number_of_heads)

    return log_posterior

x = np.linspace(1e-5, 1-1e-5, 1000)
log_posterior = compute_log_posterior(x, samples, a, b)
posterior = np.exp(log_posterior)
plt.plot(x, posterior, label='posterior', c='orange')
plt.legend()

int_posterior = 1.0 * np.mean(posterior)
print(f'Integral = {int_posterior:.4}')
```



## 5.5 Compute $\theta_{MLE}$

```
def compute_theta_mle(samples):
    """Compute theta_MLE for the given data.

    Parameters
    -----
    samples : array, shape (num_samples)
        Outcomes of simulated coin flips. Tails is 1 and heads is 0.

    Returns
    -----
    theta_mle : float
        Maximum likelihood estimate of theta.
    """
    number_of_tails = np.count_nonzero(samples==1)
    number_of_heads = len(samples) - number_of_tails

    theta_mle = number_of_tails / (number_of_tails+number_of_heads)
    return theta_mle
```

```
theta_mle = compute_theta_mle(samples)
print(f'theta_mle = {theta_mle:.3f}')
```

theta\_mle = 0.750

## 5.6 Compute $\theta_{MAP}$

```
def compute_theta_map(samples, a, b):
    """Compute theta_MAP for the given data.

    Parameters
    -----
    samples : array, shape (num_samples)
        Outcomes of simulated coin flips. Tails is 1 and heads is 0.
    a, b: float
        Parameters of the prior Beta distribution.

    Returns
    -----
    theta_map : float
        Maximum a posteriori estimate of theta.
    """
    number_of_tails = np.count_nonzero(samples==1)
    number_of_heads = len(samples) - number_of_tails

    theta_map = (number_of_tails+a+1) / (number_of_tails+2+number_of_heads+a+b)
    return theta_map
```

```
theta_map = compute_theta_map(samples, a, b)
print(f'theta_map = {theta_map:.3f}')
```

theta\_map = 0.731

## 5.7 Putting everything together

```
num_samples = 20
tails_proba = 0.7
samples = simulate_data(num_samples, tails_proba)
a, b = 5, 5
print(samples)

plt.figure(figsize=[12, 8])
x = np.linspace(1e-5, 1-1e-5, 1000)

# Plot the prior distribution
log_prior = compute_log_prior(x, a, b)
prior = np.exp(log_prior)
plt.plot(x, prior, label='prior')

# Plot the likelihood
log_likelihood = compute_log_likelihood(x, samples)
likelihood = np.exp(log_likelihood)
int_likelihood = np.mean(likelihood)
# We rescale the likelihood - otherwise it would be impossible to see in the plot
rescaled_likelihood = likelihood / int_likelihood
plt.plot(x, rescaled_likelihood, label='scaled likelihood', color='purple')

# Plot the posterior distribution
log_posterior = compute_log_posterior(x, samples, a, b)
posterior = np.exp(log_posterior)
plt.plot(x, posterior, label='posterior')

# Visualize theta_mle
theta_mle = compute_theta_mle(samples)
ymax = np.exp(compute_log_likelihood(np.array([theta_mle]), samples)) / int_likelihood
plt.vlines(x=theta_mle, ymin=0.00, ymax=ymax, linestyle='dashed', color='purple', label=r'$\theta_{MLE}$')

# Visualize theta_map
theta_map = compute_theta_map(samples, a, b)
ymax = np.exp(compute_log_posterior(np.array([theta_map]), samples, a, b))
plt.vlines(x=theta_map, ymin=0.00, ymax=ymax, linestyle='dashed', color='orange', label=r'$\theta_{MAP}$')

plt.xlabel(r'$\theta$', fontsize='xx-large')
plt.legend(fontsize='xx-large')
plt.show()
```

